

## EDAF75 Database Technology

### Lecture 3

Christian.Soderberg@cs.lth.se

January 22, 2024

## Administration

- ▶ Course representatives?
- ▶ From tomorrow you can register your lab group on the course website:
  - ▶ *make sure you enter your whole group at once, and*
  - ▶ **don't register only yourself!** (you will be removed)
- ▶ This week we'll discuss *database modeling*, and then see how you can translate a model into a database
- ▶ Lab 1 is next week, it's an exercise in SQL queries
- ▶ Lab 2 is the week after next, and it lets you practice modeling
- ▶ If you want a QA-session tomorrow, please sign up on Moodle – it's totally fine to ask question about the labs during the QA sessions

## Short recapitulation

- ▶ A *relational database* is a collection of one or more *tables*, where each table has a fixed set of *columns*, and a varying number of rows – *all cells contain primitive values*
- ▶ Simple queries (SELECT-FROM-WHERE)
- ▶ Set operations (UNION, INTERSECT, EXCEPT)
- ▶ Simple functions and aggregate functions
- ▶ Grouping (GROUP BY-HAVING)
- ▶ Window functions (OVER)
- ▶ Subqueries, views and CTEs (WITH statements)
- ▶ Joins (CROSS-, INNER-, OUTER-, ...)

## Modeling

- ▶ To design a database, we'll start out with what's called an Entity/Relation Model (E/R Model)
- ▶ There are many 'standards' for drawing E/R diagrams, we'll use UML class diagrams – it's becoming increasingly popular for database modeling

## Elements of an E/R Model

- ▶ **Entity Sets:** these are the 'objects' of our model, they correspond to classes in a traditional object oriented model
- ▶ **Attributes:** properties of our objects – must be primitive values (see the next slide)
- ▶ **Relationships:** associations between our entity sets (with cardinality)
- ▶ We will typically convert entity sets to tables (*relations*), and attributes to columns in our tables – relationships will be dealt with according to their cardinality

## 'Primitive' values in our models

- ▶ integers: INT, INTEGER, ...
- ▶ real numbers: REAL, DECIMAL(w, d), ...
- ▶ strings: TEXT, CHAR(n), VARCHAR(n)
- ▶ boolean values: true, false
- ▶ dates: DATE, TIME, TIMESTAMP, ...
- ▶ blobs (binary large objects): BLOB (only in some databases)
- ▶ JSON objects (not really primitive..., only in some databases)

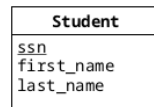
## UML class diagrams

- ▶ We'll use UML classes in approximately the same way as you may have seen them used in earlier courses, with some caveats:
  - ▶ There will be no methods in our classes
  - ▶ All our attributes will be primitive and public
  - ▶ We won't bother much with aggregates and compositions, we'll use plain associations instead
  - ▶ We'll be very careful to mark cardinalities everywhere
  - ▶ We will think carefully about what constitutes a key for each entity set

## Class diagrams for ER modeling


- ▶ We'll use:
  - ▶ classes (entity sets)
  - ▶ associations with cardinality (relationships)
  - ▶ association classes
  - ▶ inheritance (sometimes)
- ▶ We have simple rules of how to translate each kind of element into our tables
- ▶ There isn't much theory behind our ER-models, creating them is mostly an art to learn

## UML class diagrams – classes




- ▶ The entity/class name in singular
- ▶ Only one box (since we have no methods)
- ▶ We will underline keys

## UML class diagrams – associations



- ▶ We always mark cardinality on our associations
- ▶ We use associations instead of attributes whenever the value is a reference to another object


## UML class diagrams – association classes




- ▶ Sometimes the association between two entity sets contains data itself
- ▶ We use an *association class* to capture that data

## UML class diagrams – association classes

Normally we can use either an association class between two entity sets:



or another entity set 'between' them



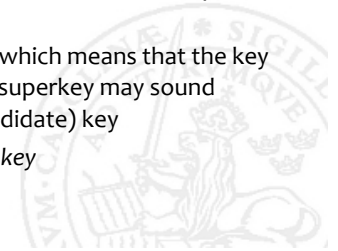
## Example

**Exercise:** Create a model for students applying for college



## Keys, candidate keys and 'super'-keys

- ▶ A *superkey* is a set of attributes for which all rows in a table are guaranteed to be unique
- ▶ A *key*, or *candidate key*, is a *minimal* set of attributes which uniquely identifies each row in a table – by minimal we mean that no attribute in the set is superfluous (it does not mean there can't be another key with fewer attributes)
- ▶ A table can have several candidate keys – when we model our database we pick one of them and call it our *primary key*
- ▶ If we add attributes to a key, the row will still be unique, which means that the key plus extra attributes is a superkey – so although being a superkey may sound impressive, it's actually less impressive than being a (candidate) key
- ▶ A key with more than one attribute is called a *composite key*




## Example

**Exercise:** What would be keys in a table of our children's classmates?



## Keys and foreign keys

- ▶ When a row in one table needs to refer to a specific row in another table (i.e., we have a  $* - 1$  association), it keeps a key to the other table (as one or more columns in this table) – this key is called a *foreign key*
- ▶ The database will ensure that there are no duplicate primary keys in a table, and it will create an *index* to speed up searches for it (more about that later), so using a primary key in another table as a foreign key makes sense
- ▶ We can also ask the database to make sure there always is a corresponding row for a foreign key (we'll return to that in a few weeks time)



## Natural keys and invented keys

- ▶ Sometimes keys occur naturally in the problem domain, we call such keys *natural keys* or *business keys*
- ▶ Entity sets which can't be uniquely identified by its attributes alone is sometimes called a *weak entity sets* (they need to use foreign keys to create a primary key) – for the sake of this course, calling them "weak" is not a big deal (it doesn't effect our designs *at all*)
- ▶ Sometimes we invent keys by introducing some artificial attribute, these keys are called *invented keys*, *surrogate keys*, *synthetic keys*, ...

## Natural keys and invented keys

- ▶ Whether to use an invented key or not is often a question of simplicity vs efficiency:
  - ▶ Without an invented key we sometimes get an unwieldy key (either because it contains many attributes, or because a single attribute might be easily mistyped)
  - ▶ With an invented key our tables and queries only need a single key column, *but finding the key may require additional joins*
- ▶ If an attribute might change over time, it's not a good choice for a key – it would require us to update all tables which uses the old value

## Weak entity sets and compound keys

- ▶ Association classes are typically 'weak', but using its foreign keys we can get a key – this is sometimes called a *compound key* (it is also a *composite key*)
- ▶ We sometimes add an invented key to a "weak" entity set – technically it's still a weak entity set (since the invented key isn't really a proper attribute)

## Example

**Exercise:** Solve the library example from the preparation web page.

## Updating or accumulating?


- ▶ How would you keep track of the balance of a bank account?
- ▶ Two ideas:
  - ▶ updating a balance attribute
  - ▶ saving all transactions, and then calculate the balance each time
- ▶ Saving the transactions allows us to track and explain the current state – it's called *Event Sourcing*, and is becoming increasingly popular
- ▶ When we update a single attribute, we need to make sure no one else updates it at the same time
- ▶ Adding a new transaction requires less locking

## Translating E/R models to databases

- ▶ For each entity set (class) we create a relation (table) with the same attributes (columns) as the entity set
- ▶ For relationships (associations), what we do depend on their cardinality
- ▶ For inheritance, there are several alternative implementations, we'll look at it next time

## Translation from model to relations – example

We have two entity sets, Book and Author, and an association between them:



Create relations for them if:

- ▶ each book is written by one author
- ▶ a book can be written by several authors

## Translation of associations

- ▶ '1 – 1'-associations typically means that the two entity sets can be merged into one entity set
- ▶ \* – 1 associations are translated into *foreign keys* on the \* side
- ▶ \* – \* associations are translated into relations with foreign keys referencing both sides
- ▶ \* – 0..1 associations are a bit special:
  - ▶ if it's mostly 1 on the right side, we can use the first method above, and use NULL where we have no associated object
  - ▶ otherwise we can use the second method above (a new relation with two foreign keys)
- ▶ Other cardinalities require some handiwork

## Translating 0..1 – 0..1 associations – example

**Exercise:** We want to keep track of people and dogs, and assume a person can only own one dog, and that a dog can be owned by at most one person.

What relations do we use if:

- ▶ Almost all dogs have an owner
- ▶ Almost every person have a dog
- ▶ Only some people own dogs, and many dogs are without an owner

## Translating 0..1 – 0..1 associations

- ▶ If almost all dogs have owners:

```
people(ssn, ...)  
dogs(id, ..., owner_ssn)
```

- ▶ If almost everyone own a dog:

```
people(ssn, ..., dog_id)  
dogs(id, ...)
```

- ▶ If only some people own dogs, and many dogs are without an owner:

```
people(ssn, ...)  
dogs(id, ...)  
dog_ownerships(owner_ssn, dog_id)
```

## Translation of association classes


- ▶ For \* – \*-relations:

- ▶ the \* – \*-relation itself will give us a new relation
- ▶ the attributes of the association class will be attributes of this relation

- ▶ For \* – 1-relations:


- ▶ the attributes of the association class will end up together with the foreign key on the \*-side

## Translating inheritance into relations



- ▶ Create one relation for each class (entity set), and reference from subclasses to superclasses using foreign keys
- ▶ Create relations only for concrete classes (entity sets)
- ▶ Create one big relation, with all possible attributes (with a lot of NULL values)

## Translating inheritance into relations




- Create one relation for each class (entity set), and reference from subclasses to superclasses using foreign keys:

```

people(ssn, name, ...)
students(ssn, stil_id, ...)
conductors(ssn, orchestra, ...)
    
```

## Translating inheritance into relations




- Create relations only for concrete classes (entity sets):

```

students(ssn, name, stil_id, ...)
conductors(ssn, name, orchestra, ...)
    
```

## Translating inheritance into relations



- Create one big relation, with all possible attributes (with a lot of NULL values)

```

people(ssn, name, stil_id, orchestra, ...)
    
```

## Defining tables in SQL

- To create a table, we use the CREATE TABLE command (and an optional DROP TABLE first):

```

DROP TABLE IF EXISTS books;
    
```

```

CREATE TABLE books (
    isbn    TEXT,
    title   TEXT,
    year    INT,
    ...
);
    
```

- For each row we define a type (see next slide)
- For each row we can add constraints



## Types in our tables

- ▶ integers: INT, INTEGER, ...
- ▶ real numbers: REAL, DECIMAL(w,d), ...
- ▶ booleans: not in SQLite, instead we use INTs, where 0 = false, and 1 = true
- ▶ strings: TEXT, CHAR(n), VARCHAR(n)
- ▶ dates: DATE, TIME, TIMESTAMP, ...
- ▶ blobs (binary large objects): BLOB (only in some databases)
- ▶ JSON objects (in some databases)

## Primary keys in table definitions

- ▶ We can declare an attribute to be primary key 'in place':

```
CREATE TABLE books (  
    isbn    TEXT PRIMARY KEY,  
    title   TEXT,  
    year    INT,  
    ...  
);
```

- ▶ We can also declare it separately (especially useful when the key contains several attributes):

```
CREATE TABLE books (  
    isbn    TEXT,  
    title   TEXT,  
    year    INT,  
    ...  
    PRIMARY KEY (isbn)  
);
```

## Foreign keys in table definitions

- ▶ We can declare foreign keys 'in place':

```
CREATE TABLE books (  
    isbn    TEXT PRIMARY KEY,  
    title   TEXT,  
    author  TEXT REFERENCES authors(author),  
    ...  
);
```

- ▶ We can also declare it separately:

```
CREATE TABLE books (  
    isbn    TEXT,  
    title   TEXT,  
    author  TEXT,  
    ...  
    FOREIGN KEY (author) REFERENCES authors(author)  
);
```

## Some other constraints

- ▶ We can declare a column to be:

- ▶ NOT NULL
- ▶ UNIQUE
- ▶ DEFAULT <value>
- ▶ CHECK <condition>

- ▶ These properties are enforced by the database, but the enforcement can often be temporarily turned off (it does take time to check everything all the time).
- ▶ We can also define *triggers* to enforce constraints, we'll return to this later in the course

## Inserting values

- ▶ We can insert values using INSERT:

```
INSERT
INTO  students(s_id, s_name, gpa, size_hs)
VALUES (123, 'Amy', 3.9, 1000),
       (234, 'Bob', 3.6, 1500),
       ...
```

- ▶ We don't have to provide values for columns with default values
- ▶ We also don't have to provide values for primary keys which are declared as INTEGER – they will get a new unique integral value (hence the moniker *database sequence number*)

## Updating values

- ▶ We can update values using UPDATE:

```
UPDATE students
SET    gpa = min(1.1 * gpa, 4.0)
WHERE  s_name LIKE 'A%';
```

- ▶ All rows are updated if we don't provide a WHERE clause

## Deleting values

- ▶ We can delete values using DELETE:

```
DELETE
FROM  applications
WHERE s_id = 123
```

- ▶ Beware that the innocent looking:

```
DELETE
FROM  applications
```

empties the whole table

## Variants

- ▶ There are various variants of the INSERT and UPDATE commands, such as:

- ▶ INSERT OR REPLACE
- ▶ INSERT OR IGNORE
- ▶ INSERT OR FAIL
- ▶ INSERT OR ROLLBACK
- ▶ UPDATE OR REPLACE
- ▶ UPDATE OR IGNORE
- ▶ UPDATE OR FAIL
- ▶ UPDATE OR ROLLBACK

- ▶ They are useful when an insertion or update would break some constraint

## Variants

- ▶ The documentation of SQLite3 gives lots of variants of the INSERT, UPDATE, and DELETE commands
- ▶ We can use WITH statements to create useful tables during inserts, updates and deletions
- ▶ We can use various forms of subqueries, etc.



## Generating invented keys

- ▶ In SQLite3 we can get a uuid-lookalike using:

```
CREATE TABLE students (  
  s_id      TEXT DEFAULT (lower(hex(randomblob(16)))),  
  s_name    TEXT,  
  gpa       DECIMAL(3,1),  
  size_hs   INT,  
  PRIMARY KEY (s_id)  
);
```

- ▶ The database doesn't have to check if the generated value is unique, since the chance of a collision is ridiculously low
- ▶ PostgreSQL has uuid as a type, SQLite3 has some library support for it, but for now it's easier to use randomblobs
- ▶ When we insert a row with generated attributes, we can have the INSERT statement return the values (using INSERT - RETURNING)

